

TEXT is faster than JSON


And JSON is faster than JSONB

* Based on some tests I did

https://github.com/waltton/pg_json_bench


TEXT vs JSON

Insert - TEXT vs JSON




Insert - TEXT vs JSON

SLOWER




Select - TEXT vs JSON




Select - TEXT vs JSON

SLOWER




TEXT vs **JSON** vs **JSONB**

Insert - TEXT vs JSON vs JSONB




Insert - TEXT vs JSON vs JSONB

SLOWER




Select - TEXT vs JSON vs JSONB



Select - TEXT vs JSON vs JSONB

SLOWER



Why?

JSON

Is basically TEXT...

○ ○ ○

```
1 test=# CREATE TABLE what_is_inside_a_json_field(data json);
2 CREATE TABLE
3
4 test=# INSERT INTO what_is_inside_a_json_field(data) VALUES ('{"postgresql": "json"}'::JSON);
5 INSERT 0 1
6
7 test=# CREATE EXTENSION pageinspect;
8 CREATE EXTENSION
9
10 test=# SELECT tuple_data_split('what_is_inside_a_json_field'::regclass, t_data, t_infomask,
    tuple_data_split
    t_infomask2, t_bits) FROM heap_page_items(get_raw_page('what_is_inside_a_json_field', 0));
11
12 -----
13 {"\\x2f7b22706f737467726573716c223a20226a736f6e227d"}
14 (1 row)
15
16 test=# \! echo '2f7b22706f737467726573716c223a20226a736f6e227d' | xxd -r -p && echo ""
17 /{"postgresql": "json"}
```

JSON

Is basically TEXT...

TEXT

○○○

```
1 test=# CREATE TABLE what_is_inside_a_json_field(data json);
2 CREATE TABLE
3
4 test=# INSERT INTO what_is_inside_a_json_field(data) VALUES ('{"postgresql": "json"}'::JSON);
5 INSERT 0 1
6
7 test=# CREATE EXTENSION pageinspect;
8 CREATE EXTENSION
9
10 test=# SELECT tuple_data_split('what_is_inside_a_json_field'::regclass, t_data, t_infomask,
    tuple_data_split
    t_infomask2, t_bits) FROM heap_page_items(get_raw_page('what_is_inside_a_json_field', 0));
11
12 -----
13 {"\\x2f7b22706f737467726573716c223a20226a736f6e227d"}
14 (1 row)
15
16 test=# \! echo '2f7b22706f737467726573716c223a20226a736f6e227d' | xxd -r -p && echo ""
17 /{"postgresql": "json"}
```

JSON

With check for JSON rules, ...

○ ○ ○

```
1 test=# CREATE TABLE tbl_text(value TEXT);
2 CREATE TABLE
3
4 test=# INSERT INTO tbl_text VALUES ('{"broken", "json'});
5 INSERT 0 1
6
7 -- -- -- -- --
8
9 test=# CREATE TABLE tbl_json(value JSON);
10 CREATE TABLE
11
12 test=# INSERT INTO tbl_json VALUES ('{"broken", "json'});
13 ERROR:  invalid input syntax for type json
14 LINE 1: INSERT INTO tbl_json VALUES ('{"broken", "json'});
15                                     ^
16 DETAIL:  Expected ":", but found ",".
17 CONTEXT:  JSON data, line 1: {"broken",...
```

JSON

With check for JSON rules, ...

INTEGRITY

○○○

```
1 test=# CREATE TABLE tbl_text(value TEXT);
2 CREATE TABLE
3
4 test=# INSERT INTO tbl_text VALUES ('{"broken", "json'});
5 INSERT 0 1
6
7 -- -- -- -- --
8
9 test=# CREATE TABLE tbl_json(value JSON);
10 CREATE TABLE
11
12 test=# INSERT INTO tbl_json VALUES ('{"broken", "json'});
13 ERROR:  invalid input syntax for type json
14 LINE 1: INSERT INTO tbl_json VALUES ('{"broken", "json'});
15                                     ^
16 DETAIL:  Expected ":", but found ", ".
17 CONTEXT:  JSON data, line 1: {"broken",...
```


JSON

And functions and operators

9.15. JSON Functions and Operators

Table 9-40 shows the operators that are available for use with JSON (see Section 8.14) data.

Table 9-40. JSON Operators

Operator	Right Operand Type	Description	Example
->	int	Get JSON array element	' [1,2,3] '::json->2
->	text	Get JSON object field	'{"a":1,"b":2}'::json->'b'
->>	int	Get JSON array element as text	' [1,2,3] '::json->>2
->>	text	Get JSON object field as text	'{"a":1,"b":2}'::json->>'b'
#>	array of text	Get JSON object at specified path	'{"a": [1,2,3], "b": [4,5,6]}'::json#>'a,2'
#>>	array of text	Get JSON object at specified path as text	'{"a": [1,2,3], "b": [4,5,6]}'::json#>>'a,2'

Table 9-41 shows the functions that are available for creating and manipulating JSON (see Section 8.14) data.

Table 9-41. JSON Support Functions

Function	Return Type	Description	Example	Example Result
array_to_json(anyarray [, pretty_bool])	json	Returns the array as JSON. A PostgreSQL multidimensional array becomes a JSON array of arrays. Line feeds will be added between dimension 1 elements if pretty_bool is true.	array_to_json('{{1,5},{99,100}}'::int[])	[[1,5],[99,100]]
row_to_json(record [, pretty_bool])	json	Returns the row as JSON. Line feeds will be added between level 1 elements if pretty_bool is true.	row_to_json(row(1,'foo'))	{"f1":1,"f2":"foo"}
to_json(anyelement)	json	Returns the value as JSON. If the data type is not built in, and there is a cast from the type to json, the cast function will be used to perform the conversion. Otherwise, for any value other than a number, a Boolean, or a null value, the text representation will be used, escaped and quoted so that it is	to_json('Fred said "Hi."'::text)	"Fred said \"Hi.\""

JSON

And functions and operators

MORE USEFUL

9.15. JSON Functions and Operators

Table 9-40 shows the operators that are available for use with JSON (see Section 8.14) data.

Table 9-40. JSON Operators

Operator	Right Operand Type	Description	Example
->	int	Get JSON array element	' [1,2,3] '::json->2
->	text	Get JSON object field	'{"a":1,"b":2}'::json->'b'
->>	int	Get JSON array element as text	' [1,2,3] '::json->>2
->>	text	Get JSON object field as text	'{"a":1,"b":2}'::json->>'b'
#>	array of text	Get JSON object at specified path	'{"a": [1,2,3], "b": [4,5,6]}'::json#>'{a,2}'
#>>	array of text	Get JSON object at specified path as text	'{"a": [1,2,3], "b": [4,5,6]}'::json#>>'{a,2}'

Table 9-41 shows the functions that are available for creating and manipulating JSON (see Section 8.14) data.

Table 9-41. JSON Support Functions

Function	Return Type	Description	Example	Example Result
array_to_json(anyarray [, pretty_bool])	json	Returns the array as JSON. A PostgreSQL multidimensional array becomes a JSON array of arrays. Line feeds will be added between dimension 1 elements if pretty_bool is true.	array_to_json('{{1,5},{99,100}}'::int[])	[[1,5],[99,100]]
row_to_json(record [, pretty_bool])	json	Returns the row as JSON. Line feeds will be added between level 1 elements if pretty_bool is true.	row_to_json(row(1,'foo'))	{"f1":1,"f2":"foo"}
to_json(anyelement)	json	Returns the value as JSON. If the data type is not built in, and there is a cast from the type to json, the cast function will be used to perform the conversion. Otherwise, for any value other than a number, a Boolean, or a null value, the text representation will be used, escaped and quoted so that it is	to_json('Fred said "Hi."'::text)	"Fred said \"Hi.\""

JSONB

B for binary, also...

```
○ ○ ○

1 test=# CREATE TABLE what_is_inside_a_jsonb_field(data jsonb);
2 CREATE TABLE
3
4 test=# INSERT INTO what_is_inside_a_jsonb_field(data) VALUES ('{"postgresql":
  "json"}'::JSONB);
5 INSERT 0 1
6
7 test=# SELECT tuple_data_split('what_is_inside_a_jsonb_field'::regclass, t_data, t_infomask,
  t_infomask2, t_bits) FROM heap_page_items(get_raw_page('what_is_inside_a_jsonb_field', 0));
8
9 tuple_data_split
10 -----
11 {"\\x37010000200a00008004000000706f737467726573716c6a736f6e"}
12
13 test=# \! echo '37010000200a00008004000000706f737467726573716c6a736f6e' | xxd -r -p && echo ""
14 7
15 postgresqljson
16 test=#
```

JSONB

B for binary, also...

BINARY

```
○ ○ ○

1 test=# CREATE TABLE what_is_inside_a_jsonb_field(data jsonb);
2 CREATE TABLE
3
4 test=# INSERT INTO what_is_inside_a_jsonb_field(data) VALUES ('{"postgresql":
  "json"}'::JSONB);
5 INSERT 0 1
6
7 test=# SELECT tuple_data_split('what_is_inside_a_jsonb_field'::regclass, t_data, t_infomask,
  t_infomask2, t_bits) FROM heap_page_items(get_raw_page('what_is_inside_a_jsonb_field', 0));
8           tuple_data_split
9  -----
10  {"\\x37010000200a00008004000000706f737467726573716c6a736f6e"}
11 (1 row)
12
13 test=# \! echo '37010000200a00008004000000706f737467726573716c6a736f6e' | xxd -r -p && echo ""
14 7
15 postgresqljson
16 test=#
```

JSONB

Faster processing

○○○

```
1 -- TEXT field
2 SELECT * FROM %s WHERE CAST(CAST(data AS JSON)->>'score' AS FLOAT) > 7.0
3
4 -- JSON and JSONB fields
5 SELECT * FROM %s WHERE CAST(data->>'score' AS FLOAT) > 7.0
```

Iterations

200


q_name

score_over_7

Table Size




Sum Duration



Duration Per Q

p50



JSONB is more than 8x faster than JSON

p90

tbl_json

p99

tbl_json

JSON is more than 1.5x faster than TEXT

JSONB

Faster processing

FASTER!!!

○○○

```
1 -- TEXT field
2 SELECT * FROM %s WHERE CAST(CAST(data AS JSON)->>'score' AS FLOAT) > 7.0
3
4 -- JSON and JSONB fields
5 SELECT * FROM %s WHERE CAST(data->>'score' AS FLOAT) > 7.0
```

Iterations

200


q_name

score_over_7

Table Size




Sum Duration



Duration Per Q

p50



JSONB is more than 8x faster than JSON

p90

tbl_json

p99

tbl_json

JSON is more than 1.5x faster than TEXT

JSONB

Indexing support

○○○

```
1 SELECT COUNT(*) FROM %s WHERE data @> '{"year": 2000}'
```

Iterations

200

q_name

count_year_2000_at_gt

Table Size



Sum Duration



▼ Duration Per Query

p50



p90

p99

Default GIN index is around 15x faster than no index on JSONB

jsonb_path_ops is almost 2x faster than the default

JSONB

Indexing support

FASTER!!!

○○○

```
1 SELECT COUNT(*) FROM %s WHERE data @> '{"year": 2000}'
```

Iterations

200

q_name

count_year_2000_at_gt

Table Size



Sum Duration



▼ Duration Per Query

p50



p90

p99

Default GIN index is around 15x faster than no index on JSONB

jsonb_path_ops is almost 2x faster than the default

Thanks



[walton/pg_json_bench](https://github.com/walton/pg_json_bench)